

Finite State Verification

Dr. Paul West

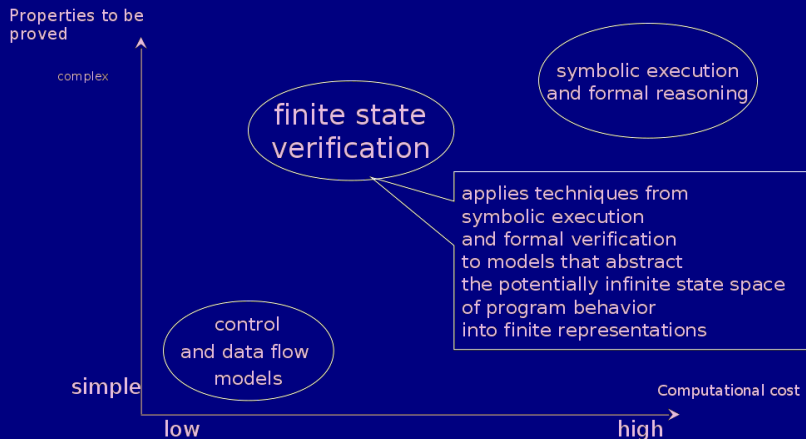
Department of Computer Science
College of Charleston

January 30, 2014

Limits and Trade-offs

- Most important properties of program execution are undecidable in general
- Finite state verification can automatically prove some significant properties of a finite model of the infinite execution space
 - balance trade-offs among
 - generality of properties to be checked
 - class of programs or models that can be checked
 - computational effort in checking
 - human effort in producing models and specifying properties

Resources and Results



Cost Trade-offs

- Human effort and skill are required
 - to prepare a finite state model
 - to prepare a suitable specification for automated analysis
- Iterative process:
 - prepare a model and specify properties
 - attempt verification
 - receive reports of impossible or unimportant faults
 - refine the specification or the model
- Automated step
 - computationally costly
 - computational cost impacts the cost of preparing model and specification, which must be tuned to make verification feasible
 - manually refining model and specification less expensive with near-interactive analysis tools

Analysis of Models

```

...
public static Table 1
getTable 1() {
  if (ref == null) {
    synchronized (Table 1) {
      if (ref == null) {
        ref = new Table 1();
        ref.initialize ();
      }
    }
  }
  return ref ;
}...

```


PROGRAM or DESIGN

Direct check of source /design
(inpractical or impossible)

PROPERTY OF INTEREST

Derive models
of software
or design

MODEL



Algorithm to check
of the model for the property

PROPERTY OF THE MODEL

never(<x> and <y>)

Implication

No concurrent
modifications of
Table 1

Application for finite State Verification

- Concurrent (multi-threaded, distributed, ...)
 - Difficult to test thoroughly (apparent non-determinism based on scheduler); sensitive to differences between development environment and field environment
 - First and most well-developed application of FSV
- Data models
 - Difficult to identify “corner cases” and interactions among constraints, or to thoroughly test them
- Security
 - Some threats depend on unusual (and untested) use

Defining the global state space - Concurrent system example

- Deriving a good finite state model is hard
- Example: finite state machine model of a program with multiple threads of control
 - Simplifying assumptions
 - we can determine in advance the number of threads
 - we can obtain a finite state machine model of each thread
 - we can identify the points at which processes can interact
 - State of the whole system model
 - = tuple of states of individual process models
 - Transition = transition of one or more of the individual processes, acting individually or in concert

State space exploration - Concurrent system example

- Specification: an on-line purchasing system
 - In-memory data structure initialized by reading configuration tables at system start-up
 - Initialization of the data structure must appear atomic
 - The system must be reinitialized on occasion
 - The structure is kept in memory
- Implementation (with bugs):
 - No monitor (Java synchronized): too expensive*
 - Double-checked locking idiom* for a fast system

*Bad decision, broken idiom ... but extremely hard to find the bug through testing.

Concurrent system example - implementation

```

class Table1 {
    private static Table1 ref = null;
    private boolean needsInit = true;
    private ElementClass [ ] theValues;
    private Table1() { }

    public static Table1 getTable1() {
        if (ref == null){
            synchedInitialize();
        }
        return ref;
    }

    private static synchronized
    void synchedInitialize() {
        if (ref == null) {
            ref = new Table1();
            ref.initialize();
        }
    }
}

```

```

public void reinit(){
    needsInit = true;
}

private synchronized void initialize() {
    . . .
    needsInit = false;
}

public int lookup(int i) {
    if (needsInit) {
        synchronized(this) {
            if (needsInit) {
                this.initialize();
            }
        }
    }
    return theValues[i].getX()
        + theValues[i].getY();
}
. . .
}

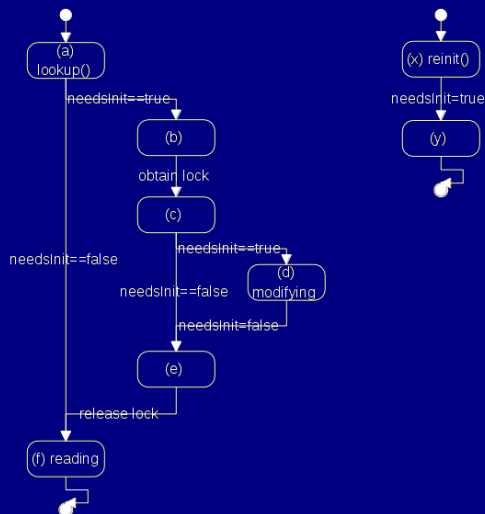
```

Analysis

- Start from models of individual threads
- Systematically trace all the possible interleavings of threads
- Like hand-executing all possible sequences of execution, but automated

... begin by constructing a finite state machine model of each individual thread ...


Finite State Machine Model for each Thread



Analysis

- Java threading rules:
 - when one thread has obtained a monitor lock
 - the other thread cannot obtain the same lock
- Locking
 - prevents threads from concurrently calling initialize
 - Does not prevent possible race condition between threads executing the lookup method
- Tracing possible executions by hand is completely impractical

Express the model in Promela



```
proctype Lookup(int id ) {  
  if :: (needsInit) ->  
    atomic { ! locked -> locked = true; };  
  if :: (needsInit) ->  
    assert (! modifying);  
    modifying = true;  
    /* Initialization happens here */  
    modifying = false ;  
    needsInit = false;  
  :: (! needsInit) ->  
    skip;  
fi;  
locked = false ;  
fi;  
assert (! modifying);}
```

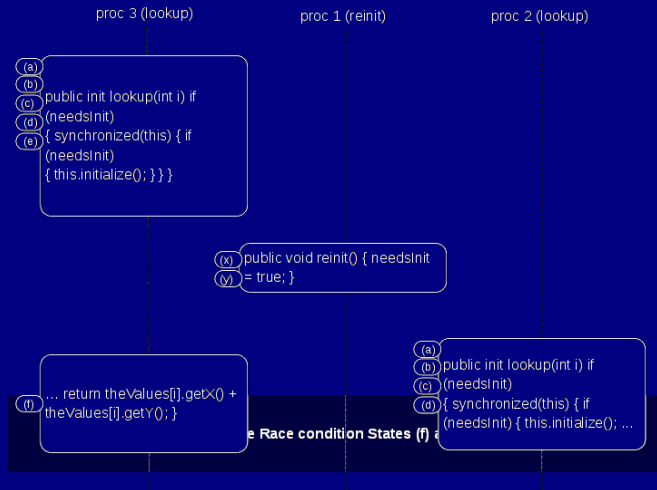
Run Spin; Inspect Output

Spin

- Depth-first search of possible executions of the model
- Explores 10 states and 51 state transitions in 0.16 seconds
- Finds a sequence of 17 transitions from the initial state of the model to a state in which one of the assertions in the model evaluates to false

```
Depth=10 States=51 Transitions=92 Memory=2.302
pan: assertion violated  !!(modifying) (at depth 17)
pan: wrote pan_in.trail
(Spin Version 4.2.5 — 2 April 2005)
0.16 real          0.00 user          0.03 sys
```

Interpret the Trace



Read/write Race condition States (f) and (d)

The Promela (Spin) Modeling Language

- A set of processes described by process types
 - Can model threads (Java), processes (Unix), devices, resources, etc.
- C-like syntax
 - expression -> statements
 - atomic statements
 - treat as a single, atomic step (without interleaving)
 - do ... od, if ... fi
 - with multiple :: alternatives, chosen non-deterministically

Safety and Liveness Properties

- Safety: bad things should not happen
 - e.g., two processes should not modify a variable at the same time.
 - Easy to specify in Promela with `assert(...)`
- Liveness: good things should eventually happen
 - e.g., if I push the button, eventually the elevator should arrive
 - Can be specified in temporal logic; more expensive to check
 - Fairness (I should get lucky now and then) is an important and common class of liveness properties

The state explosion problem

Dining philosophers - looking for deadlock with SPIN

```
5 phils+forks    145 states
                  deadlock found
10 phils+forks   18,313 states
                  error trace too long to be useful
15 phils+forks   148,897 states
                  error trace too long to be useful
```